

```
In [202]: import sqlite3
import pandas as pd
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, accuracy_score, recall_score
from math import sqrt
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.datasets import load_breast_cancer
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib import style
import seaborn as sns
%matplotlib inline
from sklearn import svm
from sklearn import datasets
import statsmodels.api as sm
#from sklearn.preprocessing import MinMaxScaler
#from sklearn.preprocessing import StandardScaler
#from sklearn.preprocessing import RobustScaler
#from sklearn.preprocessing import Normalizer
from sklearn.model_selection import cross_val_score
import statsmodels.formula.api as sm
import statsmodels.tools
```

import data

```
In [203]: df = pd.read_csv("C:/Users/mgirm/Downloads/student-mat.csv", sep=';', encoding='utf-8')
```

```
In [204]: len(df.columns)
```

```
Out[204]: 33
```

Setup a Treshold

```
In [182]: np.percentile(df.G3,[25,50,75,95,99,10])
```

```
Out[182]: array([ 8., 11., 14., 17., 19.,  5.])
```

```
In [183]: pd.Series(np.where(df.G3>14,1,0)).value_counts()
```

```
Out[183]: 0    322  
          1     73  
          dtype: int64
```

```
In [184]: df['Target']=pd.Series(np.where(df.G3>14,1,0))
```

```
In [185]: df["school"].unique()
```

```
Out[185]: array(['GP', 'MS'], dtype=object)
```

```
In [186]: df.loc[df.Target==1,:].shape[0]/df.shape[0]
```

```
Out[186]: 0.1848101265822785
```

```
In [187]: df.loc[(df.Target==1)& (df.school=='GP'),:].shape[0]/df.shape[0]
```

```
Out[187]: 0.16962025316455695
```

```
In [188]: df.loc[(df.Target==1)& (df.school=='MS'),:].shape[0]/df.shape[0]
```

```
Out[188]: 0.015189873417721518
```

of once are 18 % and when school is MS you get only 0.2%

```
In [189]: x_=df.loc[(df.Target==1)& (df.school=='GP'),:].shape[0]/df.loc[(df.school=='GP'),:].shape[0]
```

% of GP is high when value is > is 1

```
In [190]: x_/(df.loc[df.Target==1,:].shape[0]/df.shape[0])
```

```
Out[190]: 1.038780076147113
```

19 are 1

Data Dleaning and Data Exploration

In [205]: `df.isnull().any()`

```
Out[205]: school      False
sex                False
age               False
address           False
famsize           False
Pstatus           False
Medu              False
Fedu              False
Mjob              False
Fjob              False
reason            False
guardian          False
traveltime        False
studytime         False
failures          False
schoolsup         False
famsup            False
paid              False
activities        False
nursery           False
higher            False
internet          False
romantic          False
famrel            False
freetime          False
goout             False
Dalc              False
Walc              False
health            False
absences          False
G1                False
G2                False
G3                False
dtype: bool
```

In [192]: `df.shape`

Out[192]: (395, 34)

In [125]: `df.tail(5)`

```
Out[125]:
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	famrel	f
390	MS	M	20	U	LE3	A	2	2	services	services	...	5	
391	MS	M	17	U	LE3	T	3	1	services	services	...	2	
392	MS	M	21	R	GT3	T	1	1	other	other	...	5	
393	MS	M	18	R	LE3	T	3	2	services	other	...	4	
394	MS	M	19	U	LE3	T	1	1	other	at_home	...	3	

5 rows × 33 columns



In [126]: `df.columns`

Out[126]: Index(['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
 'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
 'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
 'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',
 'Walc', 'health', 'absences', 'G1', 'G2', 'G3'],
 dtype='object')

In [127]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 395 entries, 0 to 394
Data columns (total 33 columns):
school      395 non-null object
sex         395 non-null object
age         395 non-null int64
address     395 non-null object
famsize     395 non-null object
Pstatus     395 non-null object
Medu        395 non-null int64
Fedu        395 non-null int64
Mjob        395 non-null object
Fjob        395 non-null object
reason      395 non-null object
guardian    395 non-null object
traveltime  395 non-null int64
studytime   395 non-null int64
failures    395 non-null int64
schoolsup   395 non-null object
famsup      395 non-null object
paid        395 non-null object
activities  395 non-null object
nursery     395 non-null object
higher      395 non-null object
internet    395 non-null object
romantic    395 non-null object
famrel      395 non-null int64
freetime    395 non-null int64
goout       395 non-null int64
Dalc        395 non-null int64
Walc        395 non-null int64
health      395 non-null int64
absences    395 non-null int64
G1          395 non-null int64
G2          395 non-null int64
G3          395 non-null int64
dtypes: int64(16), object(17)
memory usage: 101.9+ KB
```

some students have G3 scores are 0 means student didnt attend exams hence need to drop this data

In [128]: df.describe()

Out[128]:

	age	Medu	Fedu	travelttime	studytime	failures	famrel	f
count	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395.000000	395
mean	16.696203	2.749367	2.521519	1.448101	2.035443	0.334177	3.944304	3
std	1.276043	1.094735	1.088201	0.697505	0.839240	0.743651	0.896659	0
min	15.000000	0.000000	0.000000	1.000000	1.000000	0.000000	1.000000	1
25%	16.000000	2.000000	2.000000	1.000000	1.000000	0.000000	4.000000	3
50%	17.000000	3.000000	2.000000	1.000000	2.000000	0.000000	4.000000	3
75%	18.000000	4.000000	3.000000	2.000000	2.000000	0.000000	5.000000	4
max	22.000000	4.000000	4.000000	4.000000	4.000000	3.000000	5.000000	5

In [206]: df.drop(df[df.G3==0].index,inplace=True) *#cleanup vlaues where G3 is zero i.e Target is 0*

In [130]: df.describe()

Out[130]:

	age	Medu	Fedu	travelttime	studytime	failures	famrel	f
count	357.000000	357.000000	357.000000	357.000000	357.000000	357.000000	357.000000	357
mean	16.655462	2.795518	2.546218	1.431373	2.042017	0.271709	3.955182	3
std	1.268262	1.093999	1.084217	0.686075	0.831895	0.671750	0.885721	1
min	15.000000	0.000000	0.000000	1.000000	1.000000	0.000000	1.000000	1
25%	16.000000	2.000000	2.000000	1.000000	1.000000	0.000000	4.000000	3
50%	17.000000	3.000000	3.000000	1.000000	2.000000	0.000000	4.000000	3
75%	18.000000	4.000000	3.000000	2.000000	2.000000	0.000000	5.000000	4
max	22.000000	4.000000	4.000000	4.000000	4.000000	3.000000	5.000000	5

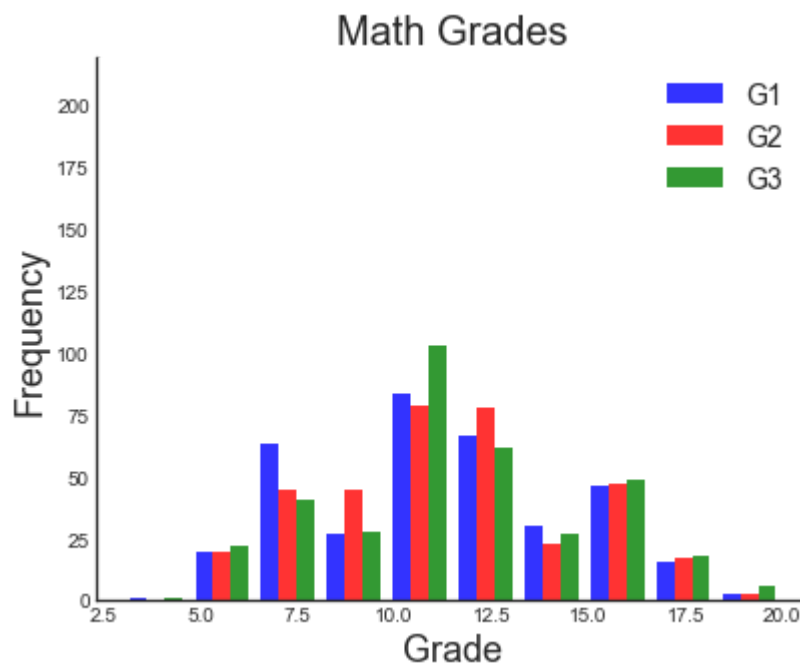
In [207]: df.groupby('sex')['G3'].mean()

Out[207]: sex
 F 11.205405
 M 11.866279
 Name: G3, dtype: float64

Plotting G1,G2 and G3 to get idea of distribution of Grades

```
In [208]: fig = plt.figure(figsize=(14,5))
plt.style.use('seaborn-white')
ax1 = plt.subplot(121)
plt.hist([df['G1'], df['G2'], df['G3']], label=['G1', 'G2', 'G3'], color=['blue', 'red', 'green'], alpha=0.8)
plt.legend(fontsize=14)
plt.xlabel('Grade', fontsize=18)
plt.ylabel('Frequency', fontsize=18)
plt.title('Math Grades', fontsize=20)
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
plt.ylim(0,220)

plt.show()
```



G1,G2 and G3 have similar distributions hence i will use G3 as final Grade to represent students performance

```
In [209]: #Math dataset
#create Aalc
df.loc[:, 'Aalc'] = (df['Dalc']*5 + df['Walc']*2)/7
#remove not interested variables
#"paid" has "no" values for all entries, so we will also drop it.
df = df.drop(['G1', 'G2', 'Dalc', 'Walc', 'paid'], axis=1)
```

In [210]: `df.columns`

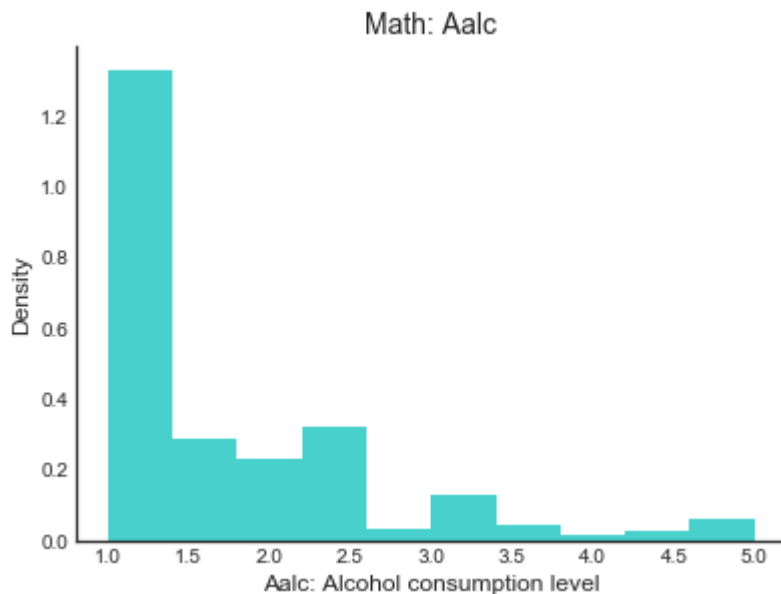
Out[210]: Index(['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
'failures', 'schoolsup', 'famsup', 'activities', 'nursery', 'higher',
'internet', 'romantic', 'famrel', 'freetime', 'goout', 'health',
'absences', 'G3', 'Aalc'],
dtype='object')

```
In [251]: #visualize Aalc
fig = plt.figure(figsize=(14,10))

ax1 = plt.subplot(221)
plt.hist(df['Aalc'], bins=10, normed=True, color='#48D1CC')
plt.title('Math: Aalc', fontsize=14)
plt.xlabel('Aalc: Alcohol consumption level', fontsize=12)
plt.ylabel('Density', fontsize=12)
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
```

C:\Users\mgirm\Anaconda3\lib\site-packages\matplotlib\axes_axes.py:6571: Use
rWarning: The 'normed' kwarg is deprecated, and has been replaced by the 'den
sity' kwarg.

warnings.warn("The 'normed' kwarg is deprecated, and has been "



From Desity plot of Aalc most stundents fall in the range of 1-1.5 meaning they consume alcholol at very low.

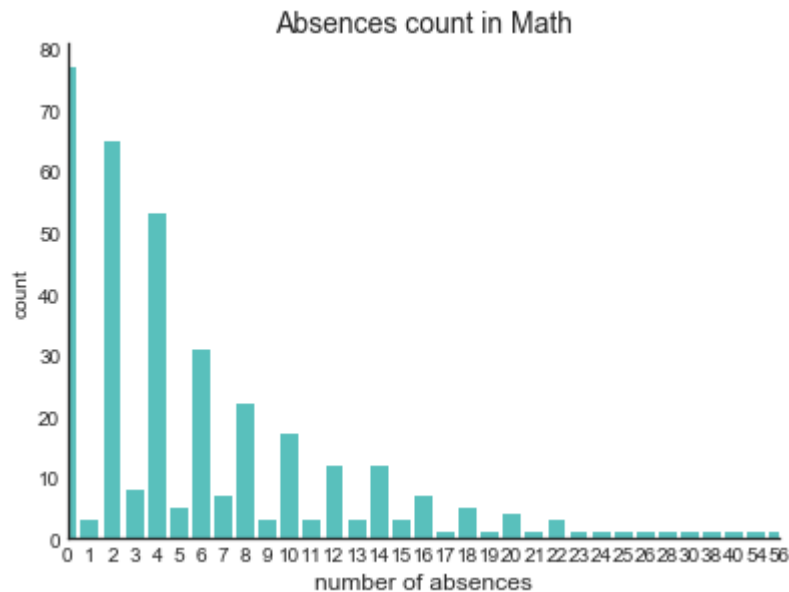
The above graph shows a reverse inverse relation between Aalc and G3

Plotting Absence

```
In [296]: #Visualize absence
fig = plt.figure(figsize=(14,10))

ax1 = plt.subplot(221)
sns.countplot(df['absences'], color='#48D1CC')
plt.title('Absences count in Math', fontsize=14)
plt.xlabel('number of absences', fontsize=12)
ax1.spines['top'].set_visible(False)
ax1.spines['right'].set_visible(False)
plt.xlim((0,32))
```

Out[296]: (0, 32)



Negative relation between G3 and Absence is shown in above plots

```
In [214]: df.tail(5)
```

Out[214]:

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	higher	i
390	MS	M	20	U	LE3	A	2	2	services	services	...	yes	
391	MS	M	17	U	LE3	T	3	1	services	services	...	yes	
392	MS	M	21	R	GT3	T	1	1	other	other	...	yes	
393	MS	M	18	R	LE3	T	3	2	services	other	...	yes	
394	MS	M	19	U	LE3	T	1	1	other	at_home	...	yes	

5 rows × 29 columns



Identifying Categorical_Columns


```
In [245]: Categorical_ = df.iloc[:,np.where([type(df[i][0])==str for i in df.columns])[0]]  
[Categorical_Columns for Categorical_Columns in Categorical_.columns]
```

```
Out[245]: ['school',  
           'sex',  
           'address',  
           'famsize',  
           'Pstatus',  
           'Mjob',  
           'Fjob',  
           'reason',  
           'guardian',  
           'schoolsup',  
           'famsup',  
           'activities',  
           'nursery',  
           'higher',  
           'internet',  
           'romantic']
```

```

In [215]: #Identify target variable y and predictor variables X.
y = df['G3']
X = df[['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
        'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
        'failures', 'schoolsup', 'famsup', 'activities', 'nursery',
        'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout',
        'health', 'absences', 'Aalc']]
#Convert dummy variables values into 0/1.
X.school = X['school'].replace(['GP', 'MS'], [1,0])
X.sex = X['sex'].replace(['F', 'M'], [1,0])
X.address = X['address'].replace(['U', 'R'], [1,0])
X.famsize = X['famsize'].replace(['LE3', 'GT3'], [1,0])
X.Pstatus = X['Pstatus'].replace(['T', 'A'], [1,0])
X.schoolsup = X['schoolsup'].replace(['yes', 'no'], [1,0])
X.famsup = X['famsup'].replace(['yes', 'no'], [1,0])
X.activities = X['activities'].replace(['yes', 'no'], [1,0])
X.nursery = X['nursery'].replace(['yes', 'no'], [1,0])
X.higher = X['higher'].replace(['yes', 'no'], [1,0])
X.internet = X['internet'].replace(['yes', 'no'], [1,0])
X.romantic = X['romantic'].replace(['yes', 'no'], [1,0])
#Identify norminal variables
norminal_vars = ['Fjob', 'Mjob', 'reason', 'guardian']
#Convert norminal variables to dummy variables
X = pd.get_dummies(X, columns = norminal_vars, drop_first=True)
# Split data into training and test data sets.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state=42)

```

C:\Users\mgirm\Anaconda3\lib\site-packages\pandas\core\generic.py:4405: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
self[name] = value

```
In [334]: # Linear Regression without GridSearch

from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn.metrics import r2_score

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3)

lm = LinearRegression()

#Next we do cross validation, which splits apart our training data and fits the model on different samples and
# gives scores for each sample to get the best fit model before we test it on the testing data.

scores = cross_val_score(lm, X_train, y_train, cv = 5)

#To get predictions (y_hat) and check them all in one using cross validation

predictions = cross_val_predict(lm, X_test, y_test, cv = 5)    #y_test is needed here in predictions to get scores for each fold of cv

accuracy = metrics.r2_score(y_test, predictions) #this says the accuracy of the predictions from the best cv fold
print("accuracy of predictions ",accuracy)

#If this is good, continue to fit the model on the data

lm.fit(X_train, y_train)

y_hat = lm.predict(X_test)    #this gives me my predictions

print("my model performance",lm.score(X_test, y_test))

accuracy of predictions  -0.02168931768172233
my model performance 0.252392653331171
```

Applying algorithms . Decision Tree


```
In [293]: from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import GridSearchCV
def decisiontree (X_train, y_train, X_test, y_test):
    param_grid = {'max_depth': range(1,100)}
    grid = GridSearchCV(DecisionTreeRegressor(), param_grid, cv=5)
    grid.fit(X_train, y_train)
    print ('Best cross validation score: {:.2f}'.format(grid.best_score_))
    print ('Best parameters:', grid.best_params_)
    print ('Test score:', grid.score(X_test, y_test))
```

In [294]: `decisiontree(X_train, y_train, X_test, y_test)`

Best cross validation score: 0.03
 Best parameters: {'max_depth': 1}
 Test score: 0.047879382104282886

```
In [295]: def plot_feature_importances_m(model):
n_features = X.shape[1]
plt.barh(range(n_features), model.feature_importances_, align='center', color='#48D1CC')
plt.yticks(np.arange(n_features), X.columns)
plt.xlabel("Feature importance")
plt.ylabel("Feature")
plt.ylim(-1, n_features)
plt.title("Math")

fig = plt.figure(figsize=(14,5))
ax1 = plt.subplot(121)
tree_m = DecisionTreeRegressor(max_depth=2).fit(X_train, y_train)
plot_feature_importances_m (tree_m)
plt.tight_layout()
```



ML Technique CV Score Test Score Parameters 1 Linear Regression 0.14 0.25 2 Decision Tree 0.22 0.15
 max_depth=2

The highest test score is 0.25, that too using Linear Regression technique. However the max score that can be achieved is 1 hence the reason for saying that the prediction models have poor performance. Using decision tree we found out that failures is the most important feature followed by schoolsup & Aalc respectively.

In []: